

EXPERIMENT 2

Implementation of Neural Network from Scratch using Numpy

OBJECTIVE

To implement a Feedforward Neural Network with backpropagation using NumPy and train it on the MNIST dataset for handwritten digit classification. The model supports multiple activation functions and includes weight initialization, gradient-based learning, and evaluation metrics.

DATA PREPROCESSING

MNIST Dataset Conversion

- The MNIST dataset was provided in binary format (`idx3-ubyte` and `idx1-ubyte`) and converted to CSV using a custom `mnist_to_csv` function.
 - Each image consists of 28x28 (784) grayscale pixels.
 - Output labels range from 0 to 9, representing digits.
 - Inputs are normalized between 0 and 1 by dividing by 255.
 - Labels are extracted and stored as integers.
-

NEURAL NETWORK IMPLEMENTATION

Architecture

- The model supports **multiple hidden layers** with custom activation functions.
- Implemented using Python classes:
 - `Layer` – stores weights, biases, and activation type.
 - `NN` – manages training, forward and backward pass, and evaluation.

```
model.add(Layer(784, 128, 'relu'))
model.add(Layer(128, 64, 'relu'))
model.add(Layer(64, 10, 'softmax'))
```

Weight Initialization

- He initialization for ReLU/Leaky ReLU.
- Xavier initialization for Sigmoid/Tanh/Softmax.

Activation Functions

| Name | Used In |
|------------|-------------------|
| ReLU | Hidden Layers |
| Sigmoid | Binary Classifier |
| Softmax | Multiclass Output |
| Tanh | Optional |
| Leaky ReLU | Optional |

TRAINING CONFIGURATION

| Hyperparameter | Value |
|----------------|---------------|
| Epochs | 1000 |
| Learning Rate | 0.01 |
| Loss Function | Cross-Entropy |
| Batch Size | Full Batch |
| Optimizer | SGD (manual) |

Training Logic

- **Forward Pass:**
 - Calculates $z = Wx + b$, followed by activation.
 - Stores intermediate activations and z values for use in backpropagation.
- **Loss Calculation:**
 - For Softmax: Categorical Cross-Entropy
 - For Sigmoid: Binary Cross-Entropy
- **Backward Pass:**
 - Computes gradients of weights and biases.
 - Updates weights using Gradient Descent.

MODEL SUMMARY


| | | |
|--------------|--------------|---------|
| Layer (type) | Output Shape | Param # |
| ===== | | |

| | | |
|---------|--------|--------|
| Layer 0 | (128,) | 100480 |
| Layer 1 | (64,) | 8256 |
| Layer 2 | (10,) | 650 |

=====

Total params: 109386


TRAINING PERFORMANCE



- Loss decreased steadily across epochs.
- Achieved stable convergence by epoch 800+.
- Final model saved to disk for later use.

```
model.save_model('./results/models/doubleLayer.pkl')
```

RESULTS & VISUALIZATION



Training Loss Plot

A graph was generated using `matplotlib`:

- Shows rapid decrease in loss early on.
- Slower convergence after epoch 500.

Test Accuracy

- Accuracy calculated by comparing predictions to ground truth.
- Class-wise accuracy charted for digits 0–9.

| Metric | Value |
|-----------------------|------------------|
| Overall Test Accuracy | 89.47% |
| Best Class Accuracy | >98% (digit '1') |
| Worst Class Accuracy | ~85% (digit '5') |

CONCLUSION

- The model generalizes well on the MNIST dataset with simple preprocessing and basic architecture.
- ReLU and Softmax activation functions yield competitive results.
- The manual implementation of backpropagation deepens the understanding of gradient flow and optimization.

FUTURE IMPROVEMENTS

- Add Dropout or L2 Regularization.
- Use Mini-Batch Gradient Descent for faster convergence.
- Extend to CNN for spatial feature learning.